

```
maria_dev@sandbox-hdp:~  
Using username "maria_dev".  
maria_dev@localhost's password:  
Last login: Wed Sep 3 13:44:52 2025 from 172.18.0.3
```

```
[maria_dev@sandbox-hdp ~]$ echo "Hadoop is big data Hadoop is java" > example.txt  
[maria_dev@sandbox-hdp ~]$ hdfs dfs -mkdir /input  
mkdir: `/input': File exists  
[maria_dev@sandbox-hdp ~]$ hdfs dfs -mkdir /Input  
[maria_dev@sandbox-hdp ~]$ hdfs dfs -put example.txt /Input/  
[maria_dev@sandbox-hdp ~]$ hdfs dfs -ls /Input  
Found 1 items  
-rw-r--r-- 1 maria_dev hdfs 34 2025-09-03 13:58 /Input/example.txt
```

```
maria_dev@sandbox-hdp:~  
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.LongWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Mapper;  
public class WordMapper extends Mapper<LongWritable, Text, Text, IntWritable> {  
    private final static IntWritable one = new IntWritable(1);  
    private Text word = new Text();  
    public void map(LongWritable key, Text value, Context context) throws IOException,  
        InterruptedException {  
        String line = value.toString();  
        for (String token : line.split("\\s+")) { // Corrected escape sequence  
            word.set(token);  
            context.write(word, one);  
        }  
    }  
}
```

```
[maria_dev@sandbox-hdp ~]$ vi WordMapper.java
```

```
maria_dev@sandbox-hdp:~  
import java.io.IOException;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Reducer;  
public class WordReducer extends Reducer<Text, IntWritable, Text, IntWritable> {  
    public void reduce(Text key, Iterable<IntWritable> values, Context context) throws IOException,  
        InterruptedException {  
        int sum = 0;  
        for (IntWritable val : values) {  
            sum += val.get();  
        }  
        context.write(key, new IntWritable(sum));  
    }  
}
```

```
[maria_dev@sandbox-hdp ~]$ vi WordReducer.java
```

```
maria_dev@sandbox-hdp:~  
import org.apache.hadoop.conf.Configuration;  
import org.apache.hadoop.fs.Path;  
import org.apache.hadoop.io.IntWritable;  
import org.apache.hadoop.io.Text;  
import org.apache.hadoop.mapreduce.Job;  
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;  
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;  
public class WordCountDriver {  
    public static void main(String[] args) throws Exception {  
        if (args.length != 2) {  
            System.err.println("Usage: WordCountDriver <input path> <output path>");  
            System.exit(-1);  
        }  
        Configuration conf = new Configuration();  
        Job job = Job.getInstance(conf, "word count");  
        job.setJarByClass(WordCountDriver.class);  
        job.setMapperClass(WordMapper.class);  
        job.setReducerClass(WordReducer.class);  
        job.setOutputKeyClass(Text.class);  
        job.setOutputValueClass(IntWritable.class);  
        FileInputFormat.addInputPath(job, new Path(args[0]));  
        FileOutputFormat.setOutputPath(job, new Path(args[1]));  
        System.exit(job.waitForCompletion(true) ? 0 : 1);  
    }  
}  
~  
~  
~  
~  
~  
~
```

```
[maria_dev@sandbox-hdp ~]$ vi WordCountDriver.java  
[maria_dev@sandbox-hdp ~]$
```

```
[maria_dev@sandbox-hdp ~]$ javac -cp `hadoop classpath` -d wordcount_classes WordMapper.java  
WordReducer.java WordCountDriver.java  
[maria_dev@sandbox-hdp ~]$ ls wordcount_classes  
WordCountDriver.class WordMapper.class WordReducer.class  
[maria_dev@sandbox-hdp ~]$
```

```
[maria_dev@sandbox-hdp ~]$ jar -cvf wordcount.jar -C wordcount_classes/ .  
added manifest  
adding: WordMapper.class(in = 1867) (out= 777) (deflated 58%)  
adding: WordReducer.class(in = 1592) (out= 664) (deflated 58%)  
adding: WordCountDriver.class(in = 1535) (out= 853) (deflated 44%)  
[maria_dev@sandbox-hdp ~]$ ls  
example.txt u.data wordcount_classes wordcount.jar WordReducer.java  
sample.txt u_local.data WordCountDriver.java WordMapper.java  
[maria_dev@sandbox-hdp ~]$
```

```

[maria_dev@sandbox-hdp:~]$ hadoop jar wordcount.jar WordCountDriver /Input /Output
25/09/03 14:11:07 INFO client.RMPProxy: Connecting to ResourceManager at sandbox-hdp.hortonworks.com/172.18.0.2:8032
25/09/03 14:11:08 INFO client.AHSPProxy: Connecting to Application History server at sandbox-hdp.hortonworks.com/172.18.0.2:10200
25/09/03 14:11:08 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
25/09/03 14:11:09 INFO input.FileInputFormat: Total input paths to process : 1
25/09/03 14:11:09 INFO mapreduce.JobSubmitter: number of splits:1
25/09/03 14:11:10 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1756907870141_0001
25/09/03 14:11:11 INFO impl.YarnClientImpl: Submitted application application_1756907870141_0001
25/09/03 14:11:12 INFO mapreduce.Job: The url to track the job: http://sandbox-hdp.hortonworks.com:8088/proxy/application_1756907870141_0001/
25/09/03 14:11:12 INFO mapreduce.Job: Running job: job_1756907870141_0001
25/09/03 14:11:36 INFO mapreduce.Job: Job job_1756907870141_0001 running in uber mode : false
25/09/03 14:11:36 INFO mapreduce.Job: map 0% reduce 0%
25/09/03 14:12:00 INFO mapreduce.Job: map 100% reduce 0%
25/09/03 14:12:25 INFO mapreduce.Job: map 100% reduce 100%
25/09/03 14:12:27 INFO mapreduce.Job: Job job_1756907870141_0001 completed successfully
25/09/03 14:12:28 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=82
    FILE: Number of bytes written=305979
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=156
    HDFS: Number of bytes written=34
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=21470
    Total time spent by all reduces in occupied slots (ms)=22039
    Total time spent by all map tasks (ms)=21470
    Total time spent by all reduce tasks (ms)=22039
    Total vcore-milliseconds taken by all map tasks=21470
    Total vcore-milliseconds taken by all reduce tasks=22039
    Total megabyte-milliseconds taken by all map tasks=5367500
    Total megabyte-milliseconds taken by all reduce tasks=5509750
  Map-Reduce Framework
    Map input records=1
    Map output records=7
    Map output materialized bytes=82
    Input split bytes=122
    Combine input records=0
    Combine output records=0
    Reduce input groups=5
    Reduce shuffle bytes=82
    Reduce input records=7
    Reduce output records=5
    Spilled Records=14
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=382
    CPU time spent (ms)=4430
    Physical memory (bytes) snapshot=342204416
    Virtual memory (bytes) snapshot=3906027520
    Total committed heap usage (bytes)=152567808
  Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
  File Input Format Counters
    Bytes Read=34
  File Output Format Counters
    Bytes Written=34
maria_dev@sandbox-hdp:~]$
```

```
[maria_dev@sandbox-hdp ~]$ hdfs dfs -ls /Output
Found 2 items
-rw-r--r--  1 maria_dev hdfs          0 2025-09-03 14:12 /Output/_SUCCESS
-rw-r--r--  1 maria_dev hdfs        34 2025-09-03 14:12 /Output/part-r-00000
[maria_dev@sandbox-hdp ~]$ hdfs dfs -cat /Output/part-r-00000
Hadoop  2
pig      1
data    1
is       2
java    1
[maria_dev@sandbox-hdp ~]$
```